

AI Incident Database: LLM-based search, assessment and review

Below are links and instructions to run an LLM-based search of the AI Incident Database (AIID) for a specified monitoring question. This requires downloading the relevant database files, downloading and customising the prompt.

Database files

- 1 Go to the AI Incident Database [download page](#) (and [homepage](#)).
- 2 Download the latest "snapshot download" (names like backup-20260914101031.tar.bz2). Copy the file name to your analysis to indicate which database version you used.
- 3 Double click to unzip the snapshot download and reveal a folder mongodump_full_snapshot that includes, among other files, incidents.csv and reports.csv.

incidents.csv

Lists, among other things, the title and description of all of the incidents in the AIID.

reports.csv

Lists all of the news reports (etc.) on which the incidents in the AIID are based. The two files are cross-referenced in incidents.csv column D.

Prompt and script

First, [download](#) the LLM prompt (aiid_assessment_prompt.md) for searching the AIID for MQ-matching incidents, and the python script (build_workbook.py) for building the assessment spreadsheet.

Open the prompt (e.g. in Visual Studio Code, another markdown editor, or simply in the LLM of your choice) and edit [PART 1](#) to add the details of your monitoring question in the spaces provided.

The prompt contains two depth/cost settings. By default, it reads every incident's description in full to judge **R** (the risk event), but does only a quick pass for **S** (the subject) on incidents where R has already been judged false — since these incidents cannot become full matches. This keeps cost down without affecting results.

You can change either setting (immediately below the monitoring question): switch S to "full" to assess S thoroughly for every incident regardless of R, or switch R to "screened" for a faster, cheaper R check that relies on a title read and keyword search — at some risk of missing a relevant incident described in a way that doesn't fit a keyword search.

Combine, run and review

- 1 Start a new chat with the LLM of your choice and select your model and effort level. The prompt was tested using Claude's Sonnet 5 model on medium effort, which is probably about right — lower effort is cheaper but increases the likelihood of skimming rather than reading.
- 2 Copy and paste in your customised prompt, and upload the `incidents.csv` file, the `reports.csv` file, and the python script `build_workbook.py`.
- 3 Alternatively, upload/open the prompt directly with your LLM prior to editing and follow its instructions to add your monitoring question details. Remember the LLM defaults to a low-effort assessment of S for any R assessed false.

The LLM should talk you through the process, query any ambiguities it detects, and explain its interpretation of your monitoring question — giving you the opportunity to modify this before beginning the analysis. Reading the prompt in full will help you understand what is happening.

At conclusion, the LLM presents results as two JSON files (recording the data efficiently) and an `xlsx` file (opened in Excel or Google Sheets) for your review. The sheet includes a "user guide" tab with detailed guidance on how to perform the review.

Human review is necessary for

- incidents deemed an indeterminate match (one or both of S and R are "indeterminate", and neither S nor R are "no");
- duplicates;
- incidents that report an ongoing rate of harm (rather than a discrete incident).

Modifications made by the user in the review sheet should automatically update the harm count in the summary tab.

The LLM should also check each matching incident for possible leads to datasets that might be of value in deriving exposure estimates.