

Incident Classification Guide

A practical aid to implementing the incident classification process — and a central repository of the resources, templates and tools that support it.

What is the incident classification process?

A process to estimate both the harm associated with an AI system and the scale of its deployment, and to describe the change in these variables between consecutive time periods in terms of five classifications: **escalating, mitigating, concentrating, receding** or **unclassifiable**. The classifications are more revealing than simple harm counts and can inform AI governance and policy.

Resources

- 01 An [interactive walkthrough](#) of the incident classification process — start here.
- 02 The AI Incident Database [homepage](#) and [download page](#).
- 03 A [digestible summary](#) of the full incident classification paper.
- 04 The [full incident classification paper](#), including five worked examples: §3.1 Case Study (p. 5) and Appendix D.1–4 (pp. 20–25).
- 05 The [working](#) behind the estimations in the full paper.

Templates

[Google Docs notes template](#)

[Google Sheets working template](#)

Tools

[Monitoring question builder](#)

[LLM assessor prompt and guide](#)

[Classifier](#)

The process in three steps

1 Define

Understand and define the type of harm, the type of victim and the cause of harm. Write these as a structured monitoring question.

2 Estimate

Search the internet for relevant data. Use it to estimate the number of incidents of harm and of exposure to harm across two consecutive time periods, then use the AI Incident Database to estimate a lower bound for the number of incidents of harm in each period.

3 Classify

Input the estimates and your best judgement of the associated uncertainty to the classifier to output the classification.

General guidance

- This document is your overall guide. The templates and tools contain more detailed instructions.
- Record your sources, your working and your thinking as if others are reading.
- Expect to iterate over and between steps 1 and 2.
- When unsure, check the worked examples for comparisons.
- Use your best judgement — this is a messy task, not a lego kit.
- Getting part way through the process and running out of data is a valid result. Document it properly. Consider adapting your monitoring question to the data you have to find a parallel result.

Notes on each step

Setting up

- Set expectations: writing and answering a monitoring question is a time-consuming, iterative, subjective, imperfect process. It can be tricky, frustrating or impossible to finish.
- Get a feel for the process with the [interactive walkthrough](#) of the Case Study (p. 5 of the [full paper](#)), the four worked examples (pp. 20–25) and the [working](#) behind all of those estimations.
- Use the templates to create a [Google Doc](#) for notes and a [Google Sheet](#) for calculations.
- This guide is not exhaustive: answering your own monitoring question will require navigating uncertainty using your best judgement.

1 Define the monitoring question

- 1.1 Begin with the [monitoring question builder](#).
- 1.2 The wording of a monitoring question affects its answer. Vague wording and vague thinking make analysis harder and answers less meaningful.
- 1.3 Check that your (S), (S+O) and (S+O+R) are countable. Be precise about what you are counting. Check that O is the specific link between S and R.
- 1.4 Read online to learn about the type of harm, the types of AI systems with which it is associated, and the specific cause of harm. This also helps identify data sources for the estimation step.
- 1.5 Try searching the [AI Incident Database](#) — incident data provides helpful context for understanding a harm type.
- 1.6 Revise and focus your monitoring question as you learn.

2 Estimate

A. Gather data

- 2.1 You need quantitative data to estimate “use” (S+O) and “harm” (S+O+R).
- 2.2 Do not expect to find data specific to (S+O) or (S+O+R). Information on any one of S, O or R may be useful in various combinations.

- 2.3 Jump ahead to section 2.13 and run the LLM-assisted search of the AIID now to identify potential data sources upfront. (You can review the full results later.)
- 2.4 All data sources in the paper were found using simple Google or LLM searches.
- 2.5 You may need to adjust your time windows to suit the data you find — and/or revise your monitoring question to accommodate available data.
- 2.6 Before quantifying anything, write a context paragraph — as if explaining to others — to organise your learning, identify gaps and record your sources.

B. Estimate harm and exposure to harm

- 2.7 The estimation process is part quantitative. The worked examples use arithmetic; fractions, ratios and percentages; [extrapolation and interpolation](#).
- 2.8 It is also part reasoning and judgement.
- 2.9 Assess each source: where does it get its data? How reliable? How complete? Can you trust it? How can you combine it with other sources? (Don't worry about methodological "tiers" at this stage — see the notes at the end.)
- 2.10 Get to your four numerical estimates however you can.
- 2.11 If you can't get to numerical estimates, can you assess a general direction (increasing or decreasing)?

C. Inspect the AIID to estimate a lower bound on incidents of harm

- 2.12 The AIID captures only a fraction of all AI harm. It forms a lower bound on actual harm.
- 2.13 Follow the [guidance](#) on how to download the AIID and use an LLM to search and assess the results.
- 2.14 This assessment is not perfect. Review the results, check high or recurring "harm count" entries, and use your judgement to interpret and adjust them.
- 2.15 There is no need for a lower bound estimate if you have a tier 1 harm estimate.
- 2.16 Practically, the lower bound is unlikely to influence the result — don't stress over it.

3 Classify

- 3.1 Enter your estimates and best judgements of uncertainty in [the classifier](#).
- 3.2 See "methodological tiers and uncertainty factors" below to help judge which tier each estimate belongs to.
- 3.3 The classifier outputs a probability for each possible classification, including "unclassifiable". Check the relative probability of each. Interpret with care.

Techniques in worked examples

Extrapolation	Case study, D1, D3
Interpolation	Case study, D1
Fractions, ratios, %	Case study, D1–D4

Intentional versus unintentional harm

Intentional harms are tricky. They often lack reliable data: attackers declare nothing, and failed attacks go unreported. All worked examples of intentional harm (D2, D3, D4) resolve as **unclassifiable**. See p. 9 of the [full paper](#).

Methodological tiers and uncertainty factors

Uncertainty in estimations comes from both the data (how reliable, how complete?) and the method (what assumptions; how is data extended or combined?). Tiers capture both.

The *uncertainty factor* (*uf*) expresses that judgement as a multiplier: for an estimate of 100 incidents, *uf* = 2 implies the “true” number lies between 50 and 200, *uf* = 3 implies 33–300, and *uf* = 1 implies the estimate is precisely accurate. As a rough rule of thumb, the tier of an estimate approximates its uncertainty factor — tier 1 → *uf* ≥ 1; tier 2 → *uf* ≥ 2; tier 3 → *uf* ≥ 3.

Tier	Data and method	<i>uf</i>
1	A single data source that is specifically relevant and which you judge relatively complete and reliable. The method requires only minor assumptions.	≥ 1
2	A combination of data sources you individually judge reliable. None is specifically relevant, but pieced together you judge them relatively complete. The method is relatively reliable, with well justified assumptions.	≥ 2
3	A combination of data sources, one or more of which you judge to be of uncertain reliability or completeness, or which require major assumptions or extrapolations in the method.	≥ 3
4	You judge the data to be insufficient, or not possible to combine in such a way as to support an estimate.	–

Straddling tiers – estimates can sit between tiers. A tier 1-to-2 estimate is more reliable than tier 2 but not quite tier 1; adjust uncertainty factors accordingly and use your best judgement. Remember: interpolation (interpreting between two known data points) is generally more reliable than extrapolation (extending data beyond the range it actually covers).